

Big Data - MongoDB



Session objectives

Big Data Overview

NoSQL introduction

MongoDB introduction

MongoDB - Java Programming



Big Data Landscape 2016

Infrastructure



Analytics



Applications



Cross-Infrastructure/Analytics



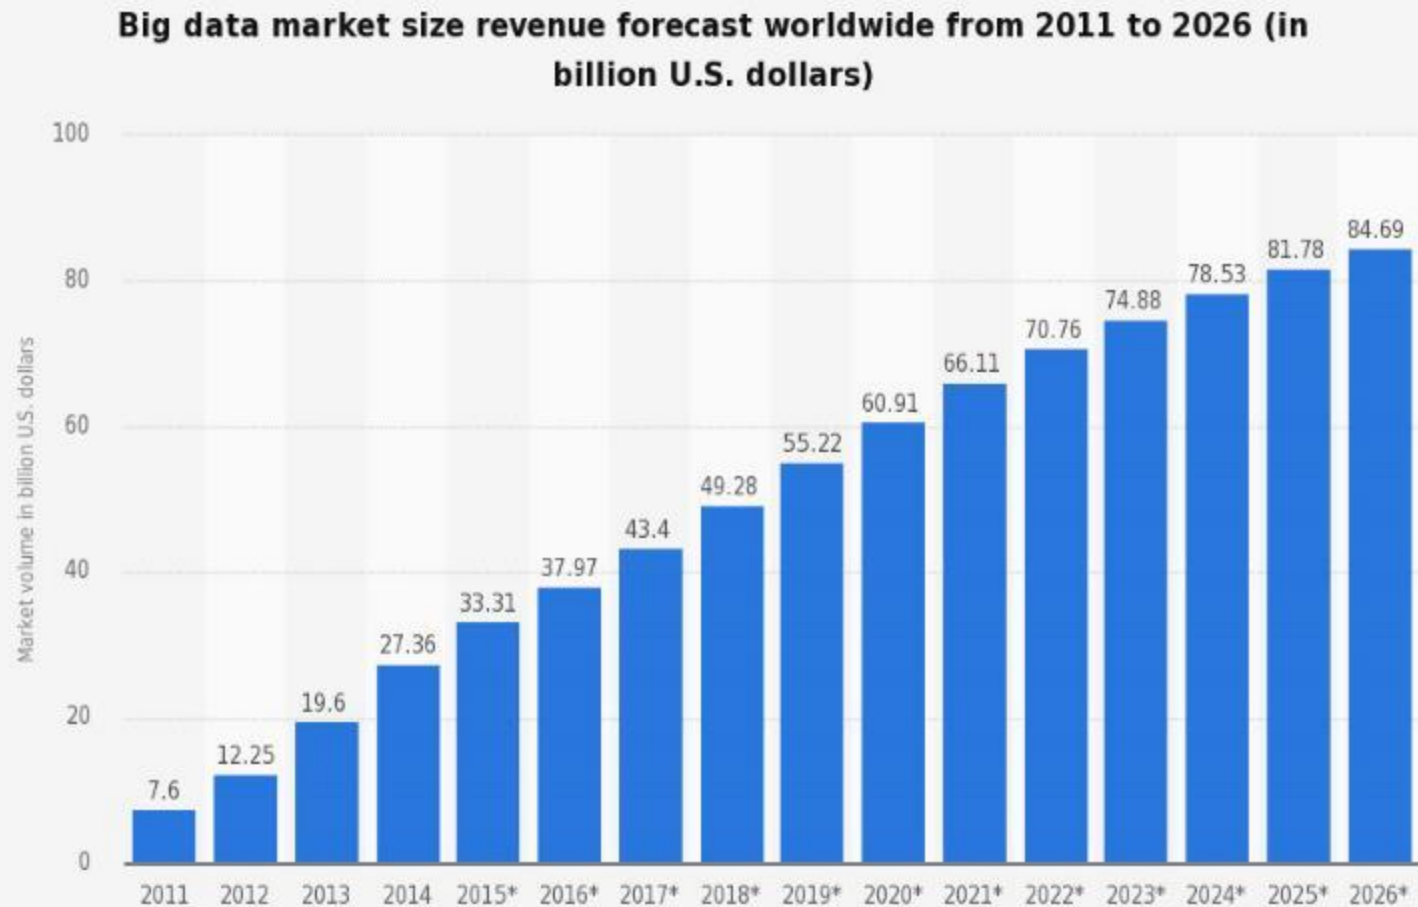
Open Source



Data Sources & APIs



Big Data, the market value



Source:
Wikibon
© Statista 2016

Additional Information:
Worldwide

Data Management Systems: History

- In the last decades RDBMS have been successful in solving problems related to storing, serving and processing data.
- RDBMS are adopted for:
 - Online transaction processing (OLTP),
 - Online analytical processing (OLAP).
- Vendors such as Oracle, Vertica, Teradata, Microsoft and IBM proposed their solution based on Relational Math and SQL.

But....




Something Changed!

- Traditionally there were transaction recording (OLTP) and analytics (OLAP) of the recorded data.
- Not much was done to understand:
 - the reasons behind transactions,
 - what factor contributed to business, and
 - what factor could drive the customer's behavior.
- Pursuing such initiatives requires working with a large amount of **varied data**.

Something Changed!

- This approach was pioneered by Google, Amazon, Yahoo, Facebook and LinkedIn.
- They work with different type of data, often semi or un-structured.
- And they have to store, serve and process huge amount of data.



Source: IDC 2011 Digital Universe Study (<http://www.emc.com/collateral/demos/microsites/emc-digital-universe-2011/index.htm>)

Something Changed!

- RDBMS can somehow deal with this aspects, but they have issues related to:
 - expensive licensing,
 - requiring complex application logic,
 - Dealing with evolving data models
- There were a need for systems that could:
 - work with different kind of data format,
 - Do not require strict schema,
 - and are easily scalable.

Evolution in Data Management

- As part of innovation in data management system, several new technologies were built:
 - 2003 - Google File System,
 - 2004 - MapReduce,
 - 2006 - BigTable,
 - 2007 - Amazon DynamoDB
 - 2012 - Google Cloud Engine
- Each solved different use cases and had a different set of assumptions.
- All these mark the beginning of a different way of thinking about data management.



Go to hell RDBMS!

Hello, Big Data!

Definition

"Big data is a term for data sets that are so large or complex that traditional data processing application software is inadequate to deal with them. Big data challenges include capturing data, data storage, data analysis, search, sharing, transfer, visualization, querying, updating and information privacy."

(https://en.wikipedia.org/wiki/Big_data)

Characteristics



- **Volume**
 - The quantity of generated and stored data. The size of the data determines the value and potential insight- and whether it can actually be considered big data or not.
- **Variety**
 - The type and nature of the data. This helps people who analyze it to effectively use the resulting insight.
- **Velocity**
 - In this context, the speed at which the data is generated and processed to meet the demands and challenges that lie in the path of growth and development.
- **Variability**
 - Inconsistency of the data set can hamper processes to handle and manage it.
- **Veracity**
 - The quality of captured data can vary greatly, affecting the accurate analysis.



NoSQL

N★SQL



NoSQL - history

- In 2006 Google published BigTable paper.
- In 2007 Amazon presented DynamoDB.
- It didn't take long for all these ideas to be used in:
 - Several open source projects (Hbase, Cassandra) and
 - Other companies (Facebook, Twitter, ...)
- And now? Now, nosql-database.org lists more than 225 NoSQL databases.



NoSQL related facts

- Explosion of social media sites (Facebook, Twitter) with large data needs.
- Rise of cloud-based solutions such as Amazon S3 (simple storage solution).
- Moving to dynamically-typed languages (Ruby/Groovy), a shift to dynamically-typed data with frequent schema changes.
- Functional Programming (Scala, Clojure, Erlang).

NoSQL Definition

<http://nosql-database.org>



"Next Generation Databases mostly addressing some of the points: being non-relational, distributed, open-source and horizontally scalable.

The original intention has been modern web-scale databases. The movement began early 2009 and is growing rapidly. Often more characteristics apply such as: schema-free, easy replication support, simple API, eventually consistent / BASE (not ACID), a huge amount of data and more. So the misleading term "nosql" (the community now translates it mostly with "not only sql") should be seen as an alias to something like the definition above."

NoSQL Categorization

1. Wide Column Store / Column Families
- 2. Document Store**
3. Key Value / Tuple Store
4. Graph Databases
5. Multimodel Databases
6. Object Databases
7. Grid & Cloud Database Solutions
8. XML Databases
9. Multidimensional Databases
10. Multivalued Databases
11. Event Sourcing
12. Time Series / Streaming Databases
13. Other NoSQL related databases
14. unresolved and uncategorized

Source: <http://nosql-database.org>

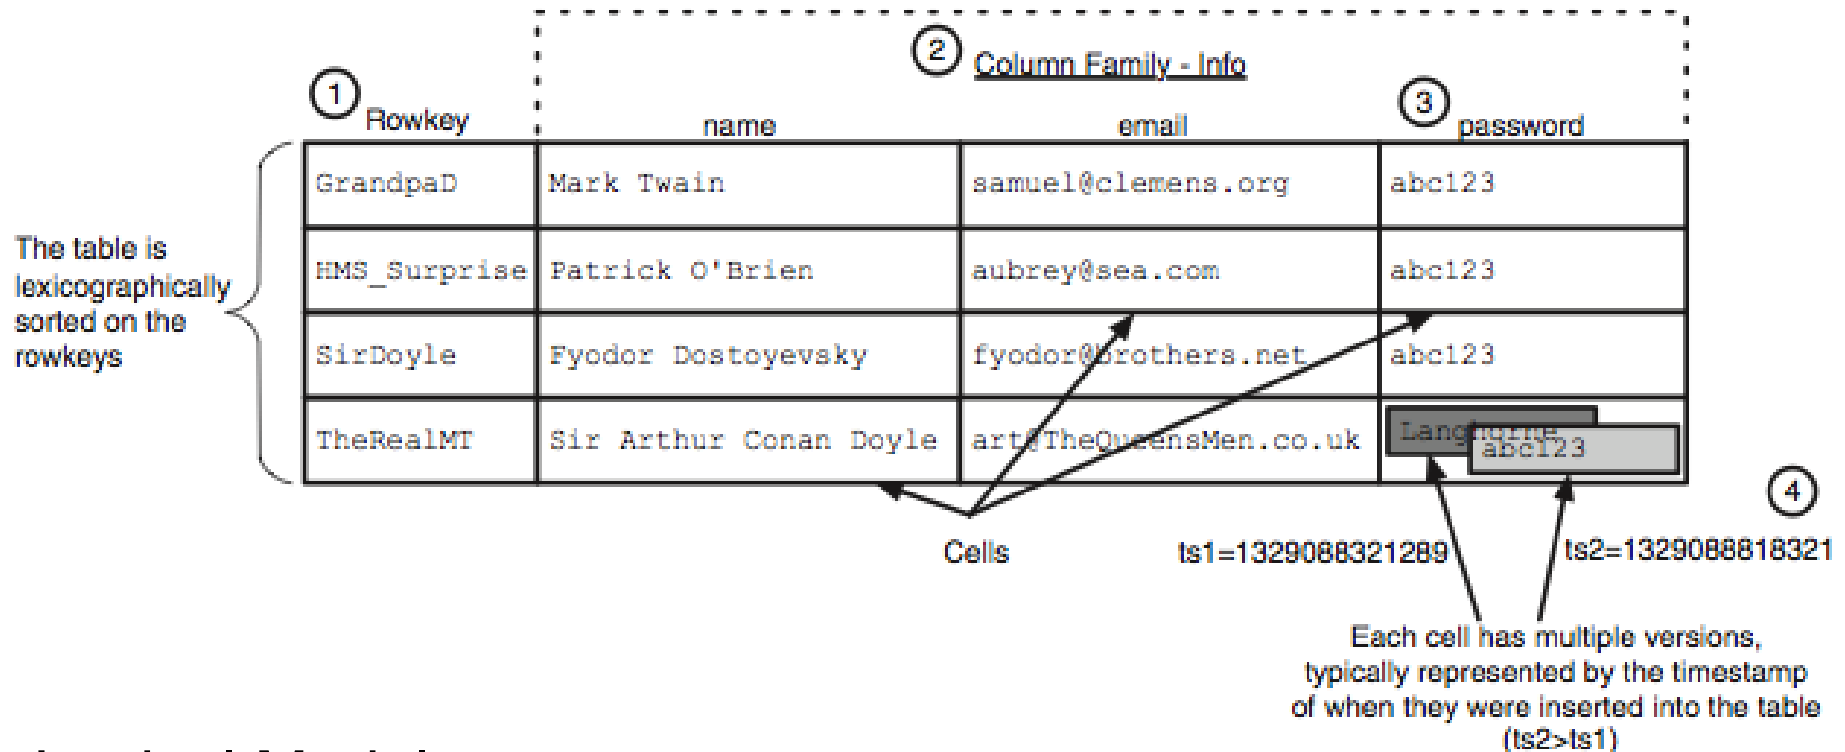
Key Value Store

- Extremely simple interface:
 - Data model: (key, value) pairs
 - Basic Operations: : Insert(key, value), Fetch(key), Update(key), Delete(key)
 - Values are store as a "blob":
 - Without caring or knowing what is inside
 - The application layer has to understand the data
 - Advantages: efficiency, scalability, fault-tolerance
- Pros:
 - very fast
 - very scalable
 - simple model
 - able to distribute horizontally
 - Cons:
 - many data structures (objects) can't be easily modeled as key value pairs

Column-oriented (1)

- Store data in columnar format
- Each storage block contains data from only one column
- Allow key-value pairs to be stored (and retrieved on key) in a massively parallel system
 - data model: families of attributes defined in a schema, new attributes can be added online
 - storing principle: big hashed distributed tables
 - properties: partitioning (horizontally and/or vertically), high availability etc. completely transparent to application

Column-oriented (2)



Logical Model

Map<RowKey, Map<ColumnFamily, Map<ColumnQualifier, Map<Version, Data>>>>

Document Store

- Schema Free.
- Usually JSON (BSON) like interchange model, which supports lists, maps, dates, Boolean with nesting
- Query Model: JavaScript or custom.
- Aggregations: Map/Reduce.
- Indexes are done via B-Trees.
- Example: Mongo
 - {Name:"Jaroslav",
Address:"Malostranske nám. 25, 118 00 Praha 1"
Grandchildren: [Claire: "7", Barbara: "6", "Magda: "3", "Kirsten: "1",
"Otis: "3", Richard: "1"]
}

Document Store: Advantages

- Documents are independent units
- Application logic is easier to write. (JSON).
- Schema Free:
 - Unstructured data can be stored easily, since a document contains whatever keys and values the application logic requires.
 - In addition, costly migrations are avoided since the database does not need to know its information schema in advance.

Graph Databases

- They are significantly different from the other three classes of NoSQL databases.
- Graph Databases are based on the mathematical concept of graph theory.
- They fit well in several real world applications (twits, permission models)
- Are based on the concepts of Vertex and Edges
- A Graph DB can be labeled, directed, attributed multi-graph
- Relational DBs can model graphs, but an edge does not require a join which is expensive.



NoSQL: How to

https://en.wikipedia.org/wiki/CAP_theorem

<http://www.julianbrowne.com/article/viewer/brewers-cap-theorem>

<https://dzone.com/articles/better-explaining-cap-theorem>



Brewer's CAP Theorem

A distributed system can support only two of the following characteristics:

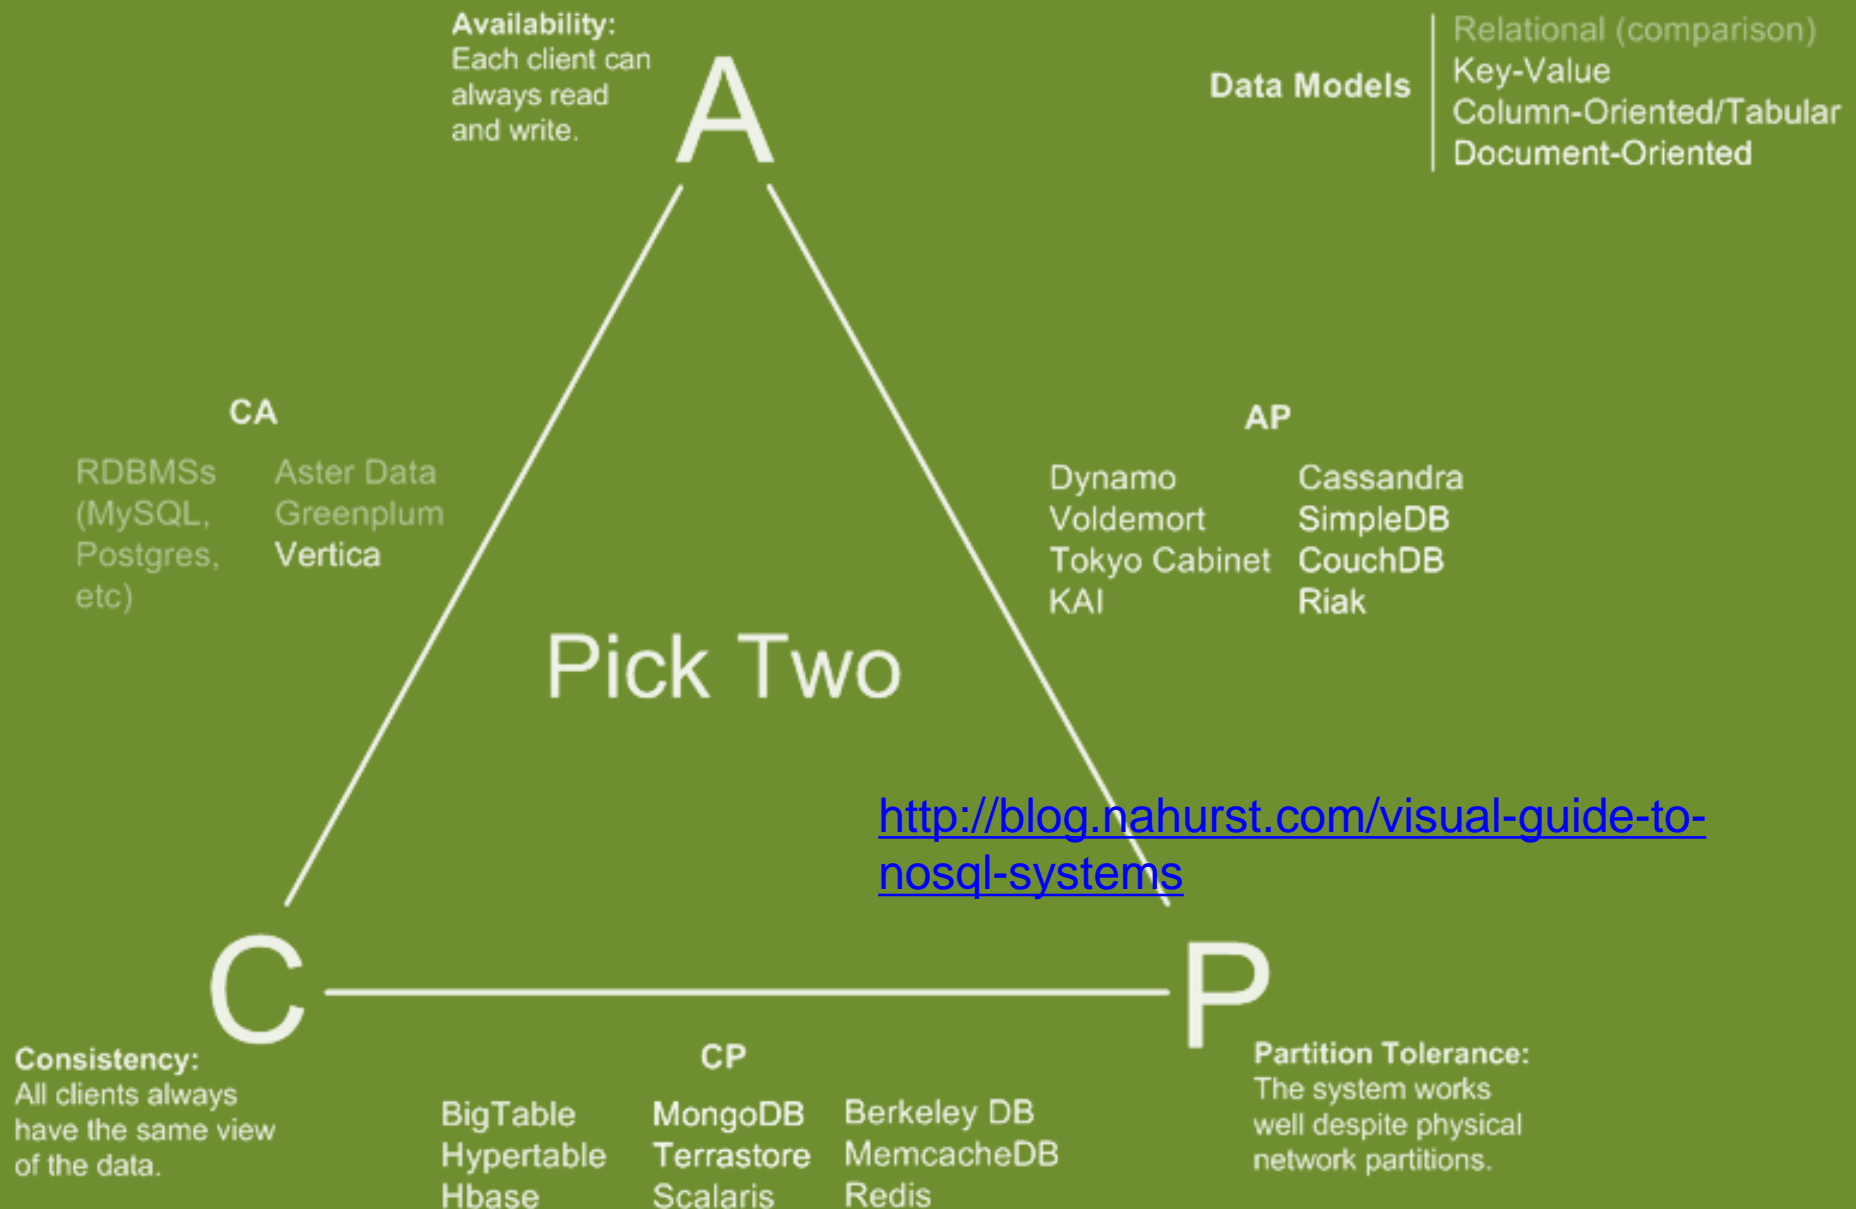
- **C**onsistency (all copies have same value)
- **A**vailability (system can run even if parts have failed)
- **P**artition Tolerance (network can break into two or more parts, each with active systems that can not influence other parts)

Brewer's CAP Theorem

Very large systems will partition at some point:

- it is necessary to decide between Consistency and Availability,
- traditional DBMS prefer Consistency over Availability and Partition,
- most Web applications choose Availability (except in specific applications such as order processing)

Visual Guide to NoSQL Systems





MongoDB




Introduction

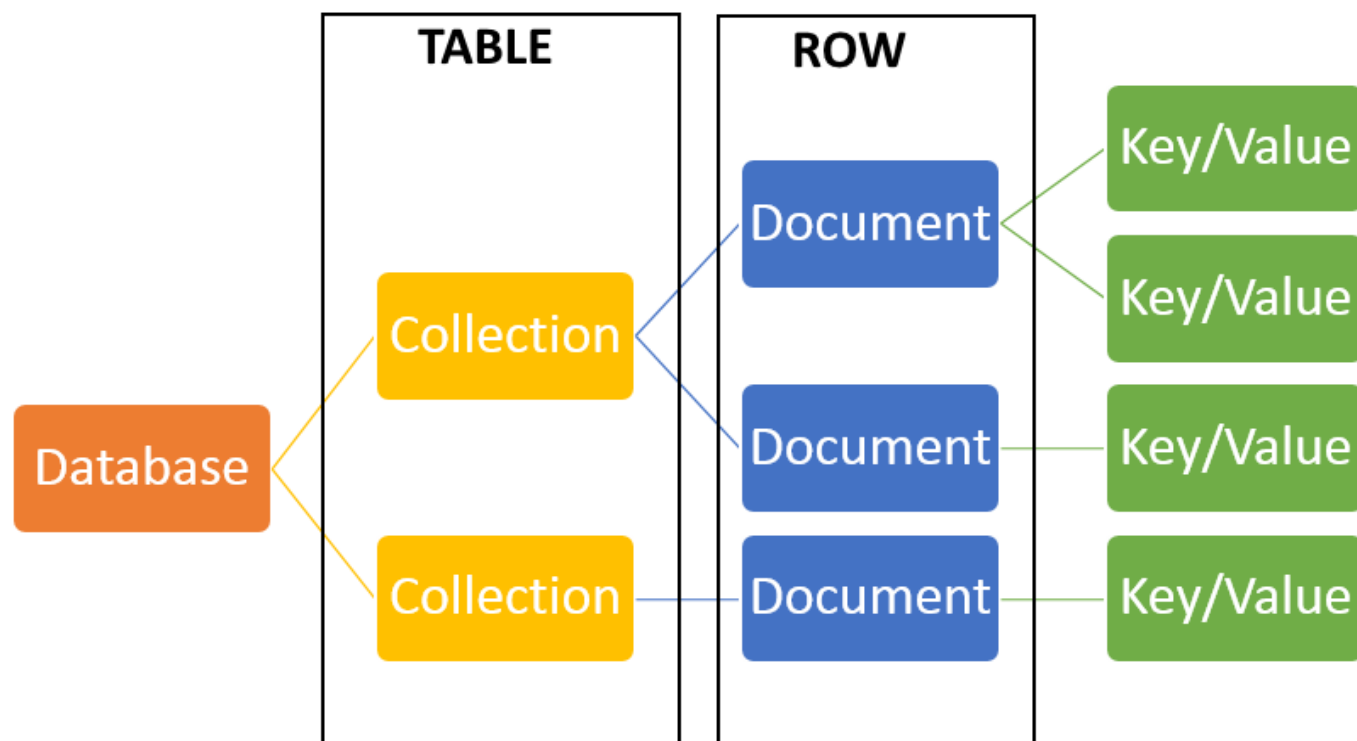
- MongoDB is an open-source database developed by MongoDB, Inc. (<https://www.mongodb.com>)
- MongoDB stores data in JSON-like (BSON) documents that can vary in structure.
- Related information is stored together for fast query access through the MongoDB query language.
- MongoDB uses dynamic schemas.

History

- 2007 - First developed (by 10gen)
- 2009 - Become Open Source
- 2010 - Considered production ready (v 1.4 >)
- 2013 - MongoDB Closes \$150 Million in Funding
- 2014 - Latest stable version (v 2.6)
- Today- More than \$231 million in total investment since 2007
- MongoDB inc. valued \$1.2B.



MongoDB structure



Terminology and Concepts

SQL Terms/Concepts	MongoDB Terms/Concepts
database	database
table	collection
row	document or BSON document
column	field
index	index
table joins	\$lookup, embedded documents
primary key Specify any unique column or column combination as primary key.	primary key In MongoDB, the primary key is automatically set to the <code>_id</code> field.
aggregation (e.g. group by)	aggregation pipeline

SQL to Aggregation Mapping Chart

SQL Terms, Functions, and Concepts	MongoDB Aggregation Operators
WHERE	<u>\$match</u>
GROUP BY	<u>\$group</u>
HAVING	<u>\$match</u>
SELECT	<u>\$project</u>
ORDER BY	<u>\$sort</u>
LIMIT	<u>\$limit</u>
SUM()	<u>\$sum</u>
COUNT()	<u>\$sum</u>
join	<u>\$lookup</u>



MongoDB - Advantages

- Flexible Data Model
- Expressive Query Syntax
- Easy to Learn
- Performance
- Scalable and Reliable
- Async Drivers
- Documentation
- Text Search
- Server-Side Script
- Documents = Objects

MongoDB - The bad

- Transactions
- No Triggers
- More Storage
- Not automatically disk cleanup
- Hierarchy of Self
- Joins
- Indexing
- Duplicate Data

Insert document

- `db.collection.insertOne()`
- `db.collection.insertMany()`

SQL INSERT Statements

```
INSERT INTO people(user_id,  
                    age,  
                    status)  
VALUES ("bcd001",  
        45,  
        "A")
```

MongoDB insertOne() Statements

```
db.people.insertOne(  
    { user_id: "bcd001", age: 45, status: "A" }  
)
```

```
try {  
    db.products.insertMany( [  
        { item: "card", qty: 15 },  
        { item: "envelope", qty: 20 },  
        { item: "stamps" , qty: 30 }  
    ] );  
} catch (e) {  
    print (e);  
}
```

Find document(s)

`db.collection.find(query, projection)`

SQL SELECT Statements

```
SELECT *  
FROM people
```

```
SELECT id,  
       user_id,  
       status  
FROM people
```

```
SELECT user_id, status  
FROM people
```

```
SELECT *  
FROM people  
WHERE status = "A"
```

MongoDB find() Statements

```
db.people.find()
```

```
db.people.find(  
  { },  
  { user_id: 1, status: 1 }  
)
```

```
db.people.find(  
  { },  
  { user_id: 1, status: 1, _id: 0 }  
)
```

```
db.people.find(  
  { status: "A" }  
)
```

```
SELECT user_id, status
FROM people
WHERE status = "A"
```

```
db.people.find(
  { status: "A" },
  { user_id: 1, status: 1, _id: 0 }
)
```

```
SELECT *
FROM people
WHERE status != "A"
```

```
db.people.find(
  { status: { $ne: "A" } }
)
```

```
SELECT *
FROM people
WHERE status = "A"
AND age = 50
```

```
db.people.find(
  { status: "A",
    age: 50 }
)
```

```
SELECT *
FROM people
WHERE status = "A"
OR age = 50
```

```
db.people.find(
  { $or: [ { status: "A" } ,
           { age: 50 } ] }
)
```

```
SELECT *  
FROM people  
WHERE age > 25
```

```
db.people.find(  
  { age: { $gt: 25 } }  
)
```

```
SELECT *  
FROM people  
WHERE age < 25
```

```
db.people.find(  
  { age: { $lt: 25 } }  
)
```

```
SELECT *  
FROM people  
WHERE age > 25  
AND age <= 50
```

```
db.people.find(  
  { age: { $gt: 25, $lte: 50 } }  
)
```

```
SELECT *  
FROM people  
WHERE user_id like "%bc%" -or-
```

```
db.people.find( { user_id: /bc/ } )
```

```
db.people.find( { user_id: { $regex: /bc/ } } )
```

```
SELECT *                                db.people.find( { user_id: /^bc/ } )
FROM people
WHERE user_id like "bc%" -or-
```

```
db.people.find( { user_id: { $regex: /^bc/ } } )
```

```
SELECT *                                db.people.find( { status: "A" } ).sort( { user_id: 1 } )
FROM people
WHERE status = "A"
ORDER BY user_id ASC
```

```
SELECT *                                db.people.find( { status: "A" } ).sort( { user_id: -1 } )
FROM people
WHERE status = "A"
ORDER BY user_id DESC
```

```
SELECT COUNT(*)                        db.people.count()
FROM people
```

or

```
db.people.find().count()
```

```
SELECT COUNT(user_id)
FROM people
```

```
db.people.count( { user_id: { $exists: true } } )
```

or

```
db.people.find( { user_id: { $exists: true } } ).count()
```

```
SELECT COUNT(*)
FROM people
WHERE age > 30
```

```
db.people.count( { age: { $gt: 30 } } )
```

or

```
db.people.find( { age: { $gt: 30 } } ).count()
```

```
SELECT DISTINCT(status)
FROM people
```

```
db.people.distinct( "status" )
```

```
SELECT *
FROM people
LIMIT 1
```

```
db.people.findOne()
```

or

```
db.people.find().limit(1)
```

```
SELECT *
FROM people
LIMIT 5
SKIP 10
```

```
db.people.find().limit(5).skip(10)
```

Explain query

```
EXPLAIN SELECT *  
FROM people  
WHERE status = "A"
```

```
db.people.find( { status: "A" } ).explain()
```

Others criteria

- limit()
- skip()
- explain()
- sort()
- count()
- pretty()
- ...

Update document

db.collection.updateOne(<filter>, <update>, <options>)

db.collection.updateMany(<filter>, <update>, <options>)

db.collection.replaceOne(<filter>, <replacement>, <options>)

SQL Update Statements

```
UPDATE people
SET status = "C"
WHERE age > 25
```

```
UPDATE people
SET age = age + 3
WHERE status = "A"
```

MongoDB updateMany() Statements

```
db.people.updateMany(
  { age: { $gt: 25 } },
  { $set: { status: "C" } }
)
```

```
db.people.updateMany(
  { status: "A" } ,
  { $inc: { age: 3 } }
)
```

Delete document

- `db.collection.deleteMany()`
- `db.collection.deleteOne()`

SQL Delete Statements

```
DELETE FROM people  
WHERE status = "D"
```

```
DELETE FROM people
```

MongoDB deleteMany() Statements

```
db.people.deleteMany( { status: "D" } )
```

```
db.people.deleteMany({})
```

Drop database

- MongoDB `db.dropDatabase()` command is used to drop a existing database.
- This will delete the selected database. If you have not selected any database, then it will delete default 'test' database.


```
>use mydb
switched to db mydb
>db.dropDatabase()
>{ "dropped" : "mydb", "ok" : 1 }
>
```

Using Management tools

MongoDB Compass

For fast schema discovery and visual construction of ad-hoc queries

- Visualize schema
 - Frequency of fields
 - Frequency of types
 - Determine validator rules
- View Documents
- Graphically build queries
- Authenticated access:



Authentication enable

- Grant permission to users to authenticate
 - Central database
 - Each database
- Policies:
 - `readAnyDatabase`
 - `readWriteAnyDatabase`
 - `userAdminAnyDatabase`
 - `dbAdminAnyDatabase`

Authentication enable

1. Create admin database
2. Add admin user

```
> use admin
switched to db admin
> db.createUser(
...  {
...    user: "admin",
...    pwd: "abc123",
...    roles: [ { role: "userAdminAnyDatabase", db: "admin" } ]
...  }
... )
Successfully added user: {
  "user" : "admin",
  "roles" : [
    {
      "role" : "userAdminAnyDatabase",
      "db" : "admin"
    }
  ]
}
```

3. Client logon:

```
mongo -u "admin" -p "abc123" -authenticationDatabase "admin"
```

- 
- Driver:

<http://mongodb.github.io/mongo-java-driver/>

- Sync

<http://mongodb.github.io/mongo-java-driver/3.5/driver/>

- A-Sync

- <http://mongodb.github.io/mongo-java-driver/3.5/driver-async/>

Connect MongoDB - sync driver

- Without authentication

```
com.mongodb.MongoClient cl=new MongoClient("localhost",27017);
```

- Authentication enable

```
List<ServerAddress>servers=new ArrayList<>();  
servers.add(new ServerAddress("localhost",27017));  
  
List<MongoCredential> credentialsList=new ArrayList<>();  
MongoCredential credential=MongoCredential.createCredential(  
    "admin", //userName  
    "admin", //authentication database  
    "abc123".toCharArray()//password  
);  
credentialsList.add(credential);  
  
com.mongodb.MongoClient mongoClient=new MongoClient(  
    servers,  
    credentialsList);
```

Get all databases

```
MongoIterable<String> ldb = mongoClient.listDatabaseNames();  
//ldb.iterator().forEachRemaining(t->{System.out.println(t);});  
  
ldb.forEach(new Block<String>() {  
    @Override  
    public void apply(String s) {  
        System.out.println(s);  
    }  
});
```

- Get specific database

```
MongoDatabase database = mongoClient.getDatabase("mondial");
```

Get collections

- Get all collections

```
MongoDatabase database = mongoClient.getDatabase("mondial");  
ListCollectionsIterable<Document> collections = database.listCollections();  
MongoIterable<String> collectionNames = database.listCollectionNames();
```

- Get specific collection

```
MongoCollection<Document> col = database.getCollection("collectionName");
```

- Create a collection

```
database.createCollection("collectionName");
```

Query

- Get all records

```
FindIterable<Document> docs = col.find();//get all  
docs.forEach(new Block<Document>() {  
    public void apply(Document t) {  
        System.out.println(t);  
    }  
});
```

- Filter criteria

```
FindIterable<Document> docs = col.find(  
    com.mongodb.client.model.Filters.eq("Name","Vietnam"));  
//--> using: import static com.mongodb.client.model.Filters.eq;
```

Insert

- Insert a Document object

```
void insert(MongoCollection<Document> col) {  
    Document lop=new Document("malop", "DHTH10A")  
        .append("tenlop", "Lớp DH Kỹ Thuật PHẦN MỀM 10A");  
    col.insertOne(lop);  
}
```

- Insert a BasicDBObject object

```
void insert2( MongoDBDatabase database) {  
    MongoCollection<BasicDBObject> collection  
        = database.getCollection("lophoc", BasicDBObject.class);  
    Map<String,String> map =new HashMap<>();  
    map.put("malop", "CDTH9LT");  
    map.put("tenlop", "Lớp cao đẳng 9 liên thông");  
    BasicDBObject bo1=new BasicDBObject(map);  
    collection.insertOne(bo1);  
}
```

Update

```
void update(MongoCollection<Document> col) {  
    //1. new data could update  
    BasicDBObject newDocument=new BasicDBObject();  
    newDocument.append("$set",  
        new BasicDBObject().append("tenlop",  
            "Lớp DHKIẾN TRÚC PHẦN MỀM 11B CLC"));  
    //2. filter object to update  
    BasicDBObject filter=new BasicDBObject().append("malop", "DHTH10A");  
  
    //3. update  
    col.updateOne(filter, newDocument);  
}
```


```
void update2(MongoCollection<Document> col) {  
    BasicDBObject filter=new BasicDBObject()  
        .append("malop", "DHTH10A");  
    BasicDBObject update=new BasicDBObject()  
        .append("$set",  
            new BasicDBObject()  
                .append("tenlop", "new value"));  
    col.findOneAndUpdate(filter, update);  
}
```

Delete

```
void delete(MongoCollection<Document> col) {  
    BasicDBObject filter=new BasicDBObject()  
        .append("malop", "1a");  
    Document doc=col.findOneAndDelete(filter);  
    //if(doc!=null) ==>success  
}
```


MongoDB - Java accessing asynchronously

- Latency
- Network traffic
- ...



Connect MongoDB - Async driver

```
MongoClient mongoClient=null;
// To directly connect to the default server localhost on port 27017
mongoClient = MongoClient.create();

/*// Use a Connection String
mongoClient = MongoClient.create("mongodb://localhost");

// or a Connection String
mongoClient = MongoClient.create(
    new ConnectionString("mongodb://localhost"));

// or provide custom MongoClientSettings
ClusterSettings clusterSettings = ClusterSettings.builder().hosts(
    Arrays.asList(new ServerAddress("localhost:27017"))
    ).build();
MongoClientSettings settings = MongoClientSettings.builder()
    .clusterSettings(clusterSettings).build();
mongoClient = MongoClient.create(settings);*/
```

Get Database & collection

```
//To directly connect to the default server localhost on port 27017
MongoClient mongoClient = MongoClient.create();

//get specific database
MongoDatabase database = mongoClient.getDatabase("mydb");

//get specific collection
MongoCollection<Document> collection = database.getCollection("test");
```

- How about the latency of network traffic?

Insert

```
//use to wait for response
final CountDownLatch latch=new CountDownLatch(1);

MongoClient mongoClient = MongoClient.create();
MongoDatabase database = mongoClient.getDatabase("mydb");
MongoCollection<Document> collection = database.getCollection("test");

Document doc = new Document("name", "MongoDB")
    .append("type", "database")
    .append("count", 1)
    .append("info", new Document("x", 203).append("y", 102));

collection.insertOne(doc, new SingleResultCallback<Void>() {
    @Override
    public void onResult(final Void result, final Throwable t) {
        System.out.println("Inserted!");
        latch.countDown();
    }
});
latch.await();
```

```
collection.insertMany(?, ?);
```

Query

```
//read all
FindIterable<Document> iter = collection.find();

iter.forEach(new Block<Document>() {
    public void apply(Document doc) {
        System.out.println(doc.toJson());
    }
}, new SingleResultCallback<Void>() {
    @Override
    public void onResult(Void v, Throwable t) {
        System.out.println("finish");
        if (t != null) {
            System.out.println("listDatabaseNames() errored: "
                + t.getMessage());
        }
        latch.countDown();
    }
});
```

Query with criteria

```
//{"i":{"$i":100}}
collection.find(gt("i", 100)).first(
    (Document doc, Throwable t)->{
        System.out.println(doc.toJson());
        latch.countDown();
    });
collection.find(gt("i", 50)).sort(ascending("pop")).forEach(
    (Document document)->{
        System.out.println(document.toJson());
    },
    (Void result, Throwable t)->{
        latch.countDown();
    }
);
```

Update

```
collection.updateOne(
    eq("i", 10), //condition
    new Document("$set", new Document("i", 110)), //update value
    new SingleResultCallback<UpdateResult>() {
        @Override
        public void onResult(final UpdateResult result,
                               final Throwable t) {
            System.out.println(result.getModifiedCount());
            latch.countDown();
        }
    });
collection.updateMany(
    lt("i", 100),
    new Document("$inc", new Document("i", 100)),
    new SingleResultCallback<UpdateResult>() {
        @Override
        public void onResult(final UpdateResult result,
                               final Throwable t) {
            System.out.println(result.getModifiedCount());
            latch.countDown();
        }
    });
```

Delete

```
collection.deleteOne(
    eq("i", 110), //condition
    new SingleResultCallback<DeleteResult>() {
        @Override
        public void onResult(final DeleteResult result, final Throwable t) {
            System.out.println(result.getDeletedCount());
            latch.countDown();
        }
    });
collection.deleteMany(
    gte("i", 100),
    new SingleResultCallback<DeleteResult>() {
        @Override
        public void onResult(final DeleteResult result, final Throwable t) {
            System.out.println(result.getDeletedCount());
            latch.countDown();
        }
    });
```


Summary

Big Data Overview

NoSQL introduction

MongoDB introduction

MongoDB - Java Programming

